

Homework 7

Problem 1

A(5, 8). B(6, 4). C(3, 5).

1.

Substitute the coordinates into the formula and we get.

$$L_1 \text{ (AC)} : \begin{bmatrix} w_1 \\ w_2 \end{bmatrix} = \begin{bmatrix} -3 \\ 2 \end{bmatrix}, \quad -1 - 3x_1 + 2x_2 = 0$$

$$L_2 \text{ (AB)} : \begin{bmatrix} w_1 \\ w_2 \end{bmatrix} = \begin{bmatrix} \frac{1}{7} \\ \frac{1}{28} \end{bmatrix}, \quad -1 + \frac{1}{7}x_1 + \frac{1}{28}x_2 = 0$$

$$L_3 \text{ (BC)} : \begin{bmatrix} w_1 \\ w_2 \end{bmatrix} = \begin{bmatrix} \frac{1}{18} \\ \frac{1}{6} \end{bmatrix}, \quad -1 + \frac{1}{18}x_1 + \frac{1}{6}x_2 = 0$$

2.

Substitute coordinate B into L_1 , we get B lies on the area. $L_1 < 0$.

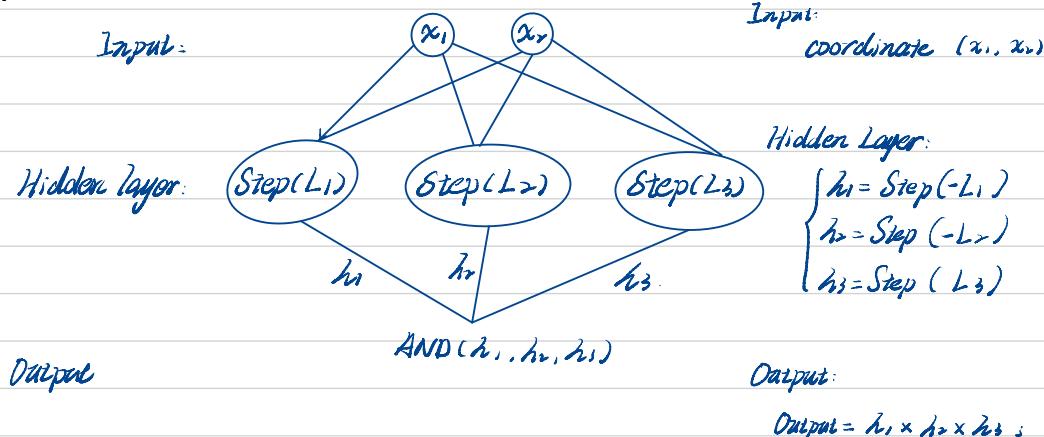
C into L_2 , C lies on the area. $L_2 < 0$.

A into L_3 , A lies on the area. $L_3 > 0$.

Therefore the 3 half planes are

$$\begin{cases} L_1 < 0 \\ L_2 < 0 \\ L_3 > 0 \end{cases}$$

3.



Problem 2

Develop a MLP to classify the quality of white wine using the dataset winequality-white.csv. It has 7 class labels (last column): 3, 4, 5, 6, 7, 8, 9. You should convert the labels into one-hot encoding. Also, you should perform feature normalization before applying them to the MLP. Partition the data (Holdout, stratified) into 60/20/20 partitions where 60% are for training, 20% are for validation and the remaining 20% are for testing (and should never be used during training)

```
In [15]: import pandas as pd
import numpy as np
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler
import torch
import torch.nn as nn
import torch.optim as optim
from torch.utils.data import DataLoader, TensorDataset
import matplotlib.pyplot as plt

# makes printing more human-friendly
np.set_printoptions(precision=3, suppress=True)
```

```
In [16]: # Import data
data = pd.read_csv('winequality-white-2.csv')

X = data.drop(columns=['quality']).values
y = data['quality'].values

y = y - 3 # Shift labels to be in range [0, 6]

# Feature standardization (Normalization)
scaler = StandardScaler()
X_normalized = scaler.fit_transform(X)

# Split data (60/20/20)
X_train, X_temp, y_train, y_temp = train_test_split(X_normalized, y, test_size=0.4, random_state=42)
X_val, X_test, y_val, y_test = train_test_split(X_temp, y_temp, test_size=0.5, random_state=42)

# Convert to PyTorch tensors
X_train = torch.tensor(X_train, dtype=torch.float32)
y_train = torch.tensor(y_train, dtype=torch.long)
X_val = torch.tensor(X_val, dtype=torch.float32)
y_val = torch.tensor(y_val, dtype=torch.long)
X_test = torch.tensor(X_test, dtype=torch.float32)
y_test = torch.tensor(y_test, dtype=torch.long)

# Create DataLoader for batch processing
batch_size = 32
train_dataset = TensorDataset(X_train, y_train)
val_dataset = TensorDataset(X_val, y_val)
test_dataset = TensorDataset(X_test, y_test)

train_loader = DataLoader(train_dataset, batch_size=batch_size, shuffle=True)
val_loader = DataLoader(val_dataset, batch_size=batch_size, shuffle=False)
test_loader = DataLoader(test_dataset, batch_size=batch_size, shuffle=False)
```

```
In [17]: # Define the MLP model
class MLP(nn.Module):
    def __init__(self, input_size, hidden_size1, hidden_size2, num_classes):
        super(MLP, self).__init__()
        self.fc1 = nn.Linear(input_size, hidden_size1)
        self.relu1 = nn.ReLU()
        self.fc2 = nn.Linear(hidden_size1, hidden_size2)
        self.relu2 = nn.ReLU()
        self.fc3 = nn.Linear(hidden_size2, num_classes)

    def forward(self, x):
        x = self.fc1(x)
        x = self.relu1(x)
        x = self.fc2(x)
        x = self.relu2(x)
        x = self.fc3(x)
        return x

# Initialize model, loss function, and optimizer
input_size = X_train.shape[1]
```

```

hidden_size1 = 64
hidden_size2 = 32
num_classes = 7

model = MLP(input_size, hidden_size1, hidden_size2, num_classes)
criterion = nn.CrossEntropyLoss()
optimizer = optim.Adam(model.parameters(), lr=0.001)

```

```

In [18]: # Training
num_epochs = 50
train_losses, val_losses = [], []

for epoch in range(num_epochs):
    model.train()
    running_loss = 0.0
    for X_batch, y_batch in train_loader:
        optimizer.zero_grad()
        outputs = model(X_batch)
        loss = criterion(outputs, y_batch)
        loss.backward()
        optimizer.step()
        running_loss += loss.item() * X_batch.size(0)
    train_loss = running_loss / len(train_loader.dataset)
    train_losses.append(train_loss)

    # Evaluate on validation set
    model.eval()
    val_loss = 0.0
    with torch.no_grad():
        for X_batch, y_batch in val_loader:
            outputs = model(X_batch)
            loss = criterion(outputs, y_batch)
            val_loss += loss.item() * X_batch.size(0)
    val_loss /= len(val_loader.dataset)
    val_losses.append(val_loss)

    print(f'Epoch [{epoch+1}/{num_epochs}], Train Loss: {train_loss:.4f}, Val Loss: {val_loss:.4f}')

# Evaluate the model on the test set
model.eval()
correct = 0
total = 0
with torch.no_grad():
    for X_batch, y_batch in test_loader:
        outputs = model(X_batch)
        _, predicted = torch.max(outputs, 1)
        total += y_batch.size(0)
        correct += (predicted == y_batch).sum().item()

test_accuracy = correct / total
print(f'Test Accuracy: {test_accuracy * 100:.2f}%')

# Visualize the training and validation losses
plt.figure(figsize=(8, 5))
plt.scatter(range(1, num_epochs + 1), train_losses, label='Training Loss')
plt.scatter(range(1, num_epochs + 1), val_losses, label='Validation Loss')
plt.xlabel('Epoch')
plt.ylabel('Loss')
plt.title('Training and Validation Loss')
plt.legend()
plt.grid()
plt.show()

```

Epoch [1/50], Train Loss: 1.4306, Val Loss: 1.1981
 Epoch [2/50], Train Loss: 1.1761, Val Loss: 1.1273
 Epoch [3/50], Train Loss: 1.1191, Val Loss: 1.1031
 Epoch [4/50], Train Loss: 1.0943, Val Loss: 1.0835
 Epoch [5/50], Train Loss: 1.0760, Val Loss: 1.0765
 Epoch [6/50], Train Loss: 1.0662, Val Loss: 1.0747
 Epoch [7/50], Train Loss: 1.0558, Val Loss: 1.0638
 Epoch [8/50], Train Loss: 1.0463, Val Loss: 1.0612
 Epoch [9/50], Train Loss: 1.0372, Val Loss: 1.0572
 Epoch [10/50], Train Loss: 1.0261, Val Loss: 1.0671
 Epoch [11/50], Train Loss: 1.0220, Val Loss: 1.0555
 Epoch [12/50], Train Loss: 1.0129, Val Loss: 1.0479
 Epoch [13/50], Train Loss: 1.0077, Val Loss: 1.0539
 Epoch [14/50], Train Loss: 1.0041, Val Loss: 1.0448
 Epoch [15/50], Train Loss: 0.9982, Val Loss: 1.0548
 Epoch [16/50], Train Loss: 0.9920, Val Loss: 1.0430
 Epoch [17/50], Train Loss: 0.9876, Val Loss: 1.0471
 Epoch [18/50], Train Loss: 0.9845, Val Loss: 1.0451
 Epoch [19/50], Train Loss: 0.9771, Val Loss: 1.0438
 Epoch [20/50], Train Loss: 0.9775, Val Loss: 1.0489
 Epoch [21/50], Train Loss: 0.9695, Val Loss: 1.0364
 Epoch [22/50], Train Loss: 0.9656, Val Loss: 1.0467
 Epoch [23/50], Train Loss: 0.9670, Val Loss: 1.0531
 Epoch [24/50], Train Loss: 0.9570, Val Loss: 1.0402
 Epoch [25/50], Train Loss: 0.9560, Val Loss: 1.0456
 Epoch [26/50], Train Loss: 0.9518, Val Loss: 1.0367
 Epoch [27/50], Train Loss: 0.9455, Val Loss: 1.0453
 Epoch [28/50], Train Loss: 0.9462, Val Loss: 1.0453
 Epoch [29/50], Train Loss: 0.9442, Val Loss: 1.0382
 Epoch [30/50], Train Loss: 0.9373, Val Loss: 1.0419
 Epoch [31/50], Train Loss: 0.9347, Val Loss: 1.0380
 Epoch [32/50], Train Loss: 0.9310, Val Loss: 1.0435
 Epoch [33/50], Train Loss: 0.9254, Val Loss: 1.0372
 Epoch [34/50], Train Loss: 0.9254, Val Loss: 1.0526
 Epoch [35/50], Train Loss: 0.9222, Val Loss: 1.0426
 Epoch [36/50], Train Loss: 0.9216, Val Loss: 1.0471
 Epoch [37/50], Train Loss: 0.9170, Val Loss: 1.0437
 Epoch [38/50], Train Loss: 0.9157, Val Loss: 1.0650
 Epoch [39/50], Train Loss: 0.9148, Val Loss: 1.0479
 Epoch [40/50], Train Loss: 0.9142, Val Loss: 1.0401
 Epoch [41/50], Train Loss: 0.9046, Val Loss: 1.0453
 Epoch [42/50], Train Loss: 0.9014, Val Loss: 1.0506
 Epoch [43/50], Train Loss: 0.9006, Val Loss: 1.0528
 Epoch [44/50], Train Loss: 0.8944, Val Loss: 1.0509
 Epoch [45/50], Train Loss: 0.8929, Val Loss: 1.0507
 Epoch [46/50], Train Loss: 0.8897, Val Loss: 1.0508
 Epoch [47/50], Train Loss: 0.8891, Val Loss: 1.0536
 Epoch [48/50], Train Loss: 0.8840, Val Loss: 1.0456
 Epoch [49/50], Train Loss: 0.8842, Val Loss: 1.0548
 Epoch [50/50], Train Loss: 0.8812, Val Loss: 1.0509
 Test Accuracy: 53.37%



Problem 3

1. AND operator

To implement AND operator. let $w_i = 1, i=1,2,\dots,n$, we have $\begin{cases} w_0 + n > 0 \\ w_0 + n - 1 < 0 \end{cases}$

Therefore $-n < w_0 < -n+1$

we have $w_0 = -n+0.5$. $w_i = 1, (i=1,2,\dots,n)$

$$w = [-n+0.5 \ 1 \ 1 \ \dots \ 1]_{n+1}$$

Only when every $x_i = 1$. $y = \text{Step}(xw) = \text{Step}(-n+0.5+n) = \text{Step}(0.5) = 1$

even only one $x_i = 0$. $y = \text{Step}(xw) = \text{Step}(-n+0.5+n-1) = \text{Step}(-0.5) = 0$

2. OR operator

To implement OR. let $w_i = 1, i=1,2,\dots,n$. we have $\begin{cases} w_0 + 0 \cdot n < 0 \\ w_0 + 1 > 0 \end{cases}$

we have $w_0 = -0.5$. $w_i = 1, (i=1,\dots,n)$

$$w = [-0.5 \ 1 \ 1 \ \dots \ 1]_{n+1}$$

Only when every x_i equals to 0. can $y = \text{Step}(xw) = \text{Step}(-0.5) = 0$.

even if only one $x_i \neq 0$. can $y = \text{Step}(xw) = \text{Step}(-0.5+1) = 1$

3. "Reverse"

Given that $M=1$, $y = \neg x_1$. we have $\begin{cases} w_0 + 1 \cdot w_1 < 0 \\ w_0 + 0 \cdot w_1 > 0 \end{cases}$

$\begin{cases} w_1 < -w_0 \\ w_0 > 0 \end{cases}$ let $w_0 = 1$. $w_1 = -2$, $w = [1 \ -2]$

when $x_1 = 1$. $y = \text{Step}(xw) = \text{Step}(1-2) = 0$

$x_1 = 0$, $y = \text{Step}(xw) = \text{Step}(1) = 1$

4. Majority vote

$$\begin{cases} w_0 + \frac{M}{2} w_1 < 0 \\ w_0 + (\frac{M}{2} + 1) w_1 > 0 \end{cases} \Rightarrow \begin{cases} w_1 > 0 \\ -\frac{M}{2} w_1 < w_0 < -\frac{M}{2} w_1 \end{cases} \text{let } w_1 = 1. w_0 = -\frac{M-1}{2}$$

when $\frac{M}{2} + 1$ $x_i = 1 \Rightarrow y = \text{Step}(-\frac{M-1}{2} + \frac{M}{2} + 1) = \text{Step}(\frac{1}{2}) = 1$

when $\frac{M}{2} x_i = 1 \Rightarrow y = \text{Step}(0) = 0$

Problem 4

```
In [19]: import random
import numpy as np
```

```
# for easier reading np
np.set_printoptions(precision=3, suppress=True)
```

```
In [20]: # Data is just a grid of points in 2D
X1, X2 = np.meshgrid(
    np.linspace(-1, 1, 20),
    np.linspace(-1, 1, 20)
)
X = np.stack((X1.flatten(), X2.flatten()), axis=1)
print('X', X.shape)

X (400, 2)
```

```
In [21]: # Part a) Specify weights and biases.
w1 = np.array([[5, 2, -5], [-5, 5, -1]])
b1 = np.array([[1, -3, 4]])
w2 = np.array([[1],[1],[1]])
b2 = np.array([[[-2.5]]])
print('w1', w1.shape) # w1 should be a 2x3 matrix
print('b1', b1.shape) # b1 should be a 1x3 matrix or a 1d vector with 3 elements
print('w2', w2.shape) # w2 should be a 3x1 matrix
print('b2', b2.shape) # b2 should be a 1x1 matrix or a 1d vector with 1 element (basically a scalar)

w1 (2, 3)
b1 (1, 3)
w2 (3, 1)
b2 (1, 1)
```

```
In [22]: # Convert to your favorite library
import torch
X = torch.from_numpy(X).float()
w1 = torch.from_numpy(w1).float()
b1 = torch.from_numpy(b1).float()
w2 = torch.from_numpy(w2).float()
b2 = torch.from_numpy(b2).float()
```

```
In [23]: # Part b) Compute the output of the network for inputs given by X.
```

```
# Hints: pytorch. Refer to:
# torch.sigmoid
# '@' implements matrix multiplication

# Compute outputs for hidden layer
h = torch.sigmoid(X @ w1 + b1)
print('h', h.shape) # h should be a 400x3 matrix

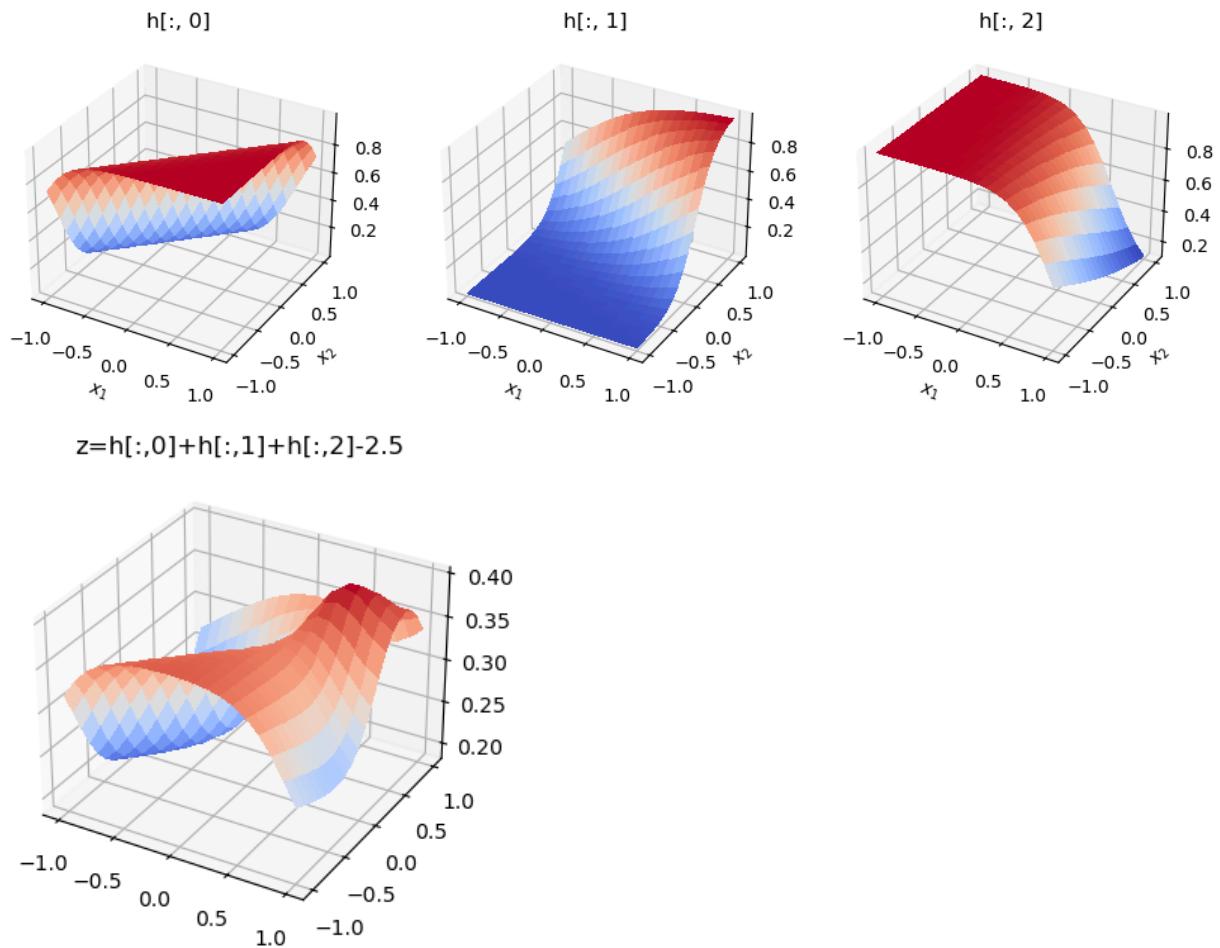
# Compute outputs for output layer
z = torch.sigmoid(h @ w2 + b2)
print('z', z.shape) # z should be a 400x1 matrix
```

```
h torch.Size([400, 3])
z torch.Size([400, 1])
```

```
In [24]: # Part c) Nothing to code.
from matplotlib import cm, pyplot as plt
fig, ax = plt.subplots(1, 3, figsize=(12, 4), subplot_kw={"projection": "3d"})
ax[0].plot_surface(X1, X2, h[:, 0].view(20, 20),
                    cmap=cm.coolwarm, linewidth=0, antialiased=False)
ax[1].plot_surface(X1, X2, h[:, 1].view(20, 20),
                    cmap=cm.coolwarm, linewidth=0, antialiased=False)
ax[2].plot_surface(X1, X2, h[:, 2].view(20, 20),
                    cmap=cm.coolwarm, linewidth=0, antialiased=False)
for i in range(3):
    ax[i].set_title(f'h[:, {i}]')
    ax[i].set_xlabel('$x_1$')
    ax[i].set_ylabel('$x_2$')

fig, ax = plt.subplots(figsize=(4, 4), subplot_kw={"projection": "3d"})
ax.plot_surface(X1, X2, z[:, 0].view(20, 20),
                cmap=cm.coolwarm, linewidth=0, antialiased=False)
ax.set_title('z=h[:,0]+h[:,1]+h[:,2]-2.5')
```

```
Out[24]: Text(0.5, 0.92, 'z=h[:,0]+h[:,1]+h[:,2]-2.5')
```



Question: How do the directions of increase of the hidden neurons relate to the weight matrix?

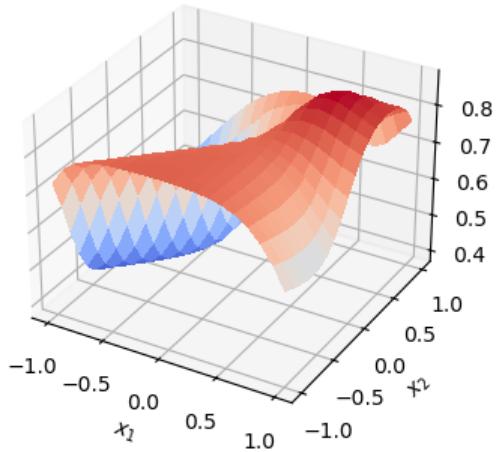
Answer: The directions of increase of the hidden neurons are the directions of increase in the weight vector.

```
In [31]: # Part d) How to increase output of the model at x=(1,-1) by changing the weights? Which weights to change, and I
w1 = torch.from_numpy(np.array([[7, 3, -5],[-7, 5, -1]]).float()
b1 = torch.from_numpy(np.array([[1, -3, 4]]).float()
w2 = torch.from_numpy(np.array([[2],[2],[2]]).float()
b2 = torch.from_numpy(np.array([-2.5])).float()

# Recompute outputs (reuse code from part b)
u1 = X @ w1 + b1
h = torch.sigmoid(u1)
u2 = h @ w2 + b2
z = torch.sigmoid(u2)

# Plot again to vizualize the results. Did the output at x=(1,-1) increased?
fig, ax = plt.subplots(figsize=(4, 4), subplot_kw={"projection": "3d"})
ax.plot_surface(X1, X2, z[:, 0].view(20, 20),
                cmap=cm.coolwarm, linewidth=0, antialiased=False)
ax.set_title('z=2*h[:,0]+h[:,1]+h[:,2]-2.5');
ax.set_xlabel('$x_1$');
ax.set_ylabel('$x_2$');
```

```
z=2*h[:,0]+h[:,1]+h[:,2]-2.5
```



```
In [26]: w1 = np.array([[5, 2, -5],[-5, 5, -1]])
b1 = np.array([[1, -3, 4]])
w2 = np.array([[1],[1],[1]])
b2 = np.array([[2.5]])
```